

Object tracking




Boris Lestsov

Goal

- 1) Build association between bboxes from frame to frame on video.

We give each detected bbox an id.



Goal

- 1) Build association between bboxes from frame to frame on video.
- 2) Fix detector errors: false positive detections and short false negatives.

We give each detected bbox an id.



Association by IoU


Given N tracks and M boxes, compute IoU matrix N.
CostMatrix is 1 - IoU.

Use hungarian algorithm to find best association between boxes and tracks. Initiate new tracks and delete old ones.

Hungarian algorithm:


N workers, M jobs => best assignment between workers and jobs in polynomial time.

IoU: 0.4034



Poor

IoU: 0.7330

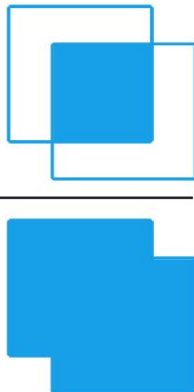


Good

IoU: 0.9264




Excellent

$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$


Association by IoU

Hungarian algorithm:

N workers, N jobs => best assignment between workers and jobs in polynomial time.



Association by IoU




Given N tracks and M boxes, compute IoU matrix N.
CostMatrix is 1 - IoU.

Use hungarian algorithm to find best association between boxes and tracks. Initiate new tracks and delete old ones.

Hungarian algorithm:

N workers, N jobs => best assignment between workers and jobs in polynomial time.

Problem: objects move on video => their boxes move => let's try to predict their movement on the next frame.

$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$


Kalman Filter

Very general and widely used in diverse applications: self-driving cars, drones, robotics, spaceships, aviation, economics, thermodynamics, ...

Data fusion: continuously fuse noisy model prediction and noisy measurements.

Kalman Filter: Outline

Object is expressed as a vector of parameters (**state vector**).


Iterate over time:

- 1) Predict state vector at the next timestep.
- 2) Correct the prediction using measurements.

Kalman filter gives MLE of parameter values (minimizes MSE in case of gaussians).

Kalman Filter: Example

One dimensional cart that can move along one axis.




Kalman Filter: Model

Express information about object as a normal distribution.

P_k - **covariance matrix**, our uncertainty about state vector.

In example: position and velocity can be correlated => matrix is not diagonal.

$$\hat{\mathbf{x}}_k = \begin{bmatrix} \text{position} \\ \text{velocity} \end{bmatrix}$$
$$\mathbf{P}_k = \begin{bmatrix} \Sigma_{pp} & \Sigma_{pv} \\ \Sigma_{vp} & \Sigma_{vv} \end{bmatrix}$$



Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

Kalman Filter: Model

Express state transformation with **prediction matrix**.
Depends on environment.

The modelled process is linear.

Our example: cart moves with some speed => its position changes.

Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

$$p_k = p_{k-1} + \Delta t v_{k-1}$$

$$v_k = v_{k-1}$$

$$\begin{aligned}\hat{\mathbf{x}}_k &= \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix} \hat{\mathbf{x}}_{k-1} \\ &= \mathbf{F}_k \hat{\mathbf{x}}_{k-1}\end{aligned}$$

$$\begin{aligned}\text{Cov}(x) &= \Sigma \\ \text{Cov}(\mathbf{A}x) &= \mathbf{A}\Sigma\mathbf{A}^T\end{aligned}$$

$$\begin{aligned}\hat{\mathbf{x}}_k &= \mathbf{F}_k \hat{\mathbf{x}}_{k-1} \\ \mathbf{P}_k &= \mathbf{F}_k \mathbf{P}_{k-1} \mathbf{F}_k^T\end{aligned}$$

Kalman Filter: External Control

u_k - **control vector**.

B_k - **control matrix**.

Our example: an external force is applied to the cart.

$$p_k = p_{k-1} + \Delta t v_{k-1} + \frac{1}{2} a \Delta t^2$$
$$v_k = v_{k-1} + a \Delta t$$

$$\hat{\mathbf{x}}_k = \mathbf{F}_k \hat{\mathbf{x}}_{k-1} + \begin{bmatrix} \frac{\Delta t^2}{2} \\ \Delta t \end{bmatrix} a$$
$$= \mathbf{F}_k \hat{\mathbf{x}}_{k-1} + \mathbf{B}_k \vec{u}_k$$

Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

Kalman Filter: External Noise


The environment might itself might cause uncertainty =>
add matrix \mathbf{Q}_k - **model noise covariance matrix**.

Our example: friction adds noise.

$$\begin{aligned}\hat{\mathbf{x}}_k &= \mathbf{F}_k \hat{\mathbf{x}}_{k-1} + \mathbf{B}_k \vec{\mathbf{u}}_k \\ \mathbf{P}_k &= \mathbf{F}_k \mathbf{P}_{k-1} \mathbf{F}_k^T + \mathbf{Q}_k\end{aligned}$$


Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

Kalman Filter: Measurement



Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

Kalman Filter: Measurement



Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

Kalman Filter: Measurement


Measurement and state vectors may have different size, because not all state vector elements may be observable.

In trivial case, H is diagonal.

$$\begin{aligned}\vec{\mu}_{\text{expected}} &= \mathbf{H}_k \hat{\mathbf{x}}_k \\ \Sigma_{\text{expected}} &= \mathbf{H}_k \mathbf{P}_k \mathbf{H}_k^T\end{aligned}$$

Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

Kalman Filter: Measurement



Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>


Kalman Filter: Combining estimates

We have two estimates:

- 1) Predicted by the model
- 2) Measured


Both are gaussians.

How do we combine them?




Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

Kalman Filter: Combining estimates



Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

Kalman Filter: Combining estimates



$$\mathcal{N}(x, \mu_0, \sigma_0) \cdot \mathcal{N}(x, \mu_1, \sigma_1) \stackrel{?}{=} \mathcal{N}(x, \mu', \sigma')$$

Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

Kalman Filter: Combining estimates

$$\mathbf{K} = \Sigma_0(\Sigma_0 + \Sigma_1)^{-1}$$

$$\vec{\mu}' = \vec{\mu}_0 + \mathbf{K}(\vec{\mu}_1 - \vec{\mu}_0)$$

$$\Sigma' = \Sigma_0 - \mathbf{K}\Sigma_0$$

K - Kalman Gain

Source: <https://www.bzarg.com/p/how-a-kalman-filter-works-in-pictures/>

Kalman Filter: Update

The predicted measurement with $(\mu_0, \Sigma_0) = (\mathbf{H}_k \hat{\mathbf{x}}_k, \mathbf{H}_k \mathbf{P}_k \mathbf{H}_k^T)$,
and the observed measurement with $(\mu_1, \Sigma_1) = (\vec{\mathbf{z}}_k, \mathbf{R}_k)$.

$$\hat{\mathbf{x}}'_k = \hat{\mathbf{x}}_k + \mathbf{K}'(\vec{\mathbf{z}}_k - \mathbf{H}_k \hat{\mathbf{x}}_k)$$

$$\mathbf{P}'_k = \mathbf{P}_k - \mathbf{K}' \mathbf{H}_k \mathbf{P}_k$$

$$\mathbf{K}' = \mathbf{P}_k \mathbf{H}_k^T (\mathbf{H}_k \mathbf{P}_k \mathbf{H}_k^T + \mathbf{R}_k)^{-1} \quad - \quad \text{Kalman Gain}$$

Kalman Filter

Kalman Filter hyperparameters:

F - state transition matrix.

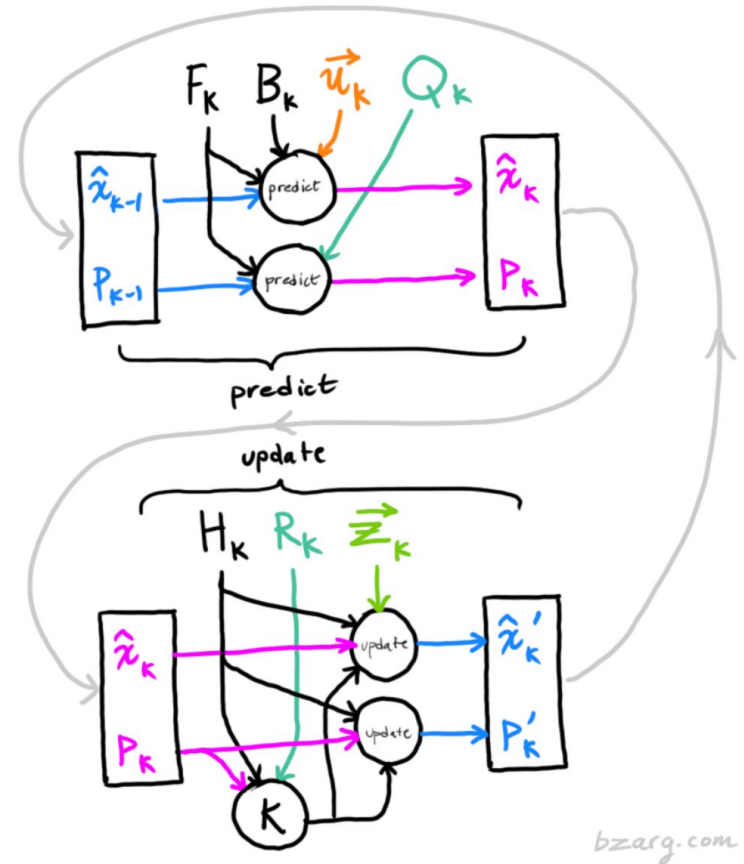
H - measurement matrix.

P - initial state uncertainty covariance matrix

Q - model noise covariance (model uncertainty).

R - measurement noise covariance (measurement uncertainty).

Kalman Filter Information Flow



Kalman Filter: Questions



Kalman Filter: Assignment

Notebook: seminar1-kalman.ipynb



Zip archive in Discord



Work in groups

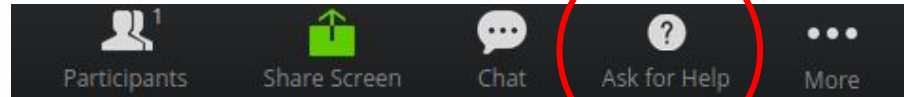


40 minutes



Ask your teammates

Call the lecturer if you need help



SORT: Questions



Kalman Filter: Questions




SORT

SORT: states and measurements

Each tracklet = Kalman Filter with state $[x, y, ar, h, x', y', h']$:

1. Bbox center on x axis.
2. Bbox center on y axis.
3. Aspect ratio.
4. Height.
5. Movement speed along x.
6. Movement speed along y.
7. Height change speed.



$$ar = h/w$$

Detections are used as measurements for Kalman update: state vector size is 7, measurement vector size is 4 (only $[x, y, ar, h]$).

SORT: assignment

How to associate new detections to tracks?

SORT: assignment

How to associate new detections to tracks? - **by IoU**.

Linear assignment problem on matrix:

CostMatrix = Tracks x Detections.

Hungarian algorithm is used.

SORT

Minimum IoU threshold of predicted bbox and tracked bbox is required.

If IoU for some detection is less than threshold => new track is initiated.

SORT: parameters

We keep “age” of each track (frame count).

- 1) **max_age** - How many frames to wait until track is deleted.
- 2) **min_hits** - Minimal track age required to keep the track.
- 3) **iou_thresh** - Minimal IoU threshold to match detection with a track.

If track is finished before **min_hits** => it is ignored (detector False Positive).

Track states:

- 1) Initiated: found unmatched detection
- 2) Confirmed : age > **min_hits**
- 3) Missed : no detection at this timestep
- 4) Deleted : time since last update > **max_age**

SORT: usage, tips, applications

- 1) Works very fast on CPU.
- 2) Minimum FPS is required: about 5 FPS for faces/pedestrians.
- 3) Tune tracker hyperparameters (previous slide).

Applications:

- 1) Face/Object Recognition:
 - a) Compute centroid embedding for “good” subset of crops in a track => AKNN => propagate label on full track. Much better embedding quality.
 - b) Keep only 1 centroid for each track in database. 1 embedding instead of embeddings of all objects in a track => faster AKNN and smaller base.
 - c) Video Analytics.

SORT: Questions



SORT: Assignment

Notebook: seminar2-sort.ipynb



Zip archive in Discord



Work in groups




30 minutes



Ask your teammates

Call the lecturer if you need help



DeepSORT

- 1) Applicable if we have object descriptors (embeddings).
- 2) Associate detections with tracklets by embedding similarity.
- 3) Unmatched detections are associated by IoU (covered faces/pedestrians).
- 4) Matching cascade.
- 5) Gating mechanism.

DeepSORT: Matching cascade

Keep a collection of embeddings from N previous frames.

At timestep T , try to assign detections to tracks with hungarian algorithm. Start with embeddings from step $T-1$, $T-2$, etc..

Use MSE/Cosine Distance between embeddings to compute CostMatrix.

DeepSORT: Matching cascade

Keep a collection of embeddings from N previous frames.

At timestep T , try to assign detections to tracks with hungarian algorithm. Start with embeddings from step $T-1$, $T-2$, etc..

Use MSE/Cosine Distance between embeddings to compute CostMatrix.

In the end, try to match all unmatched detections with IoU CostMatrix.

Problem: We do not use any spatial information about tracks and bboxes. => use gating mechanism.

DeepSORT: Gating

Modify CostMatrix with gating values:

If distance between i-th measurement and j-th state (it's projection) is greater than threshold \Rightarrow CostMatrix_{ij} is replaced with infinity/bigger value.



Track state is a normal **distribution** \Rightarrow we can measure distance from point to mean, but it's bad.

Problem: how to measure a distance from point to distribution? \Rightarrow mahalanobis distance

DeepSORT: Gating

Modify CostMatrix with gating values:

If mahalanobis distance between i-th measurement and j-th state (it's projection) is greater than threshold => CostMatrix_{ij} is replaced with infinity/bigger value.



DeepSORT: Mahalanobis Distance

Squared Mahalanobis distance: $d^{(1)}(i, j) = (\mathbf{d}_j - \mathbf{y}_i)^T \mathbf{S}_i^{-1} (\mathbf{d}_j - \mathbf{y}_i),$

$\mathbf{y}_i, \mathbf{S}_i$ - i-th track state projection into measurement space

\mathbf{d}_j - j-th detection.

“Distance from point to distribution”.

Relation to MLE:

$\ln(L) = -d^2 - \text{const}$

DeepSORT: Overview

At timestep T:

- Update kalman states.

- For t in $[T-1, T-2, \dots, T-N]$:

 - Compute CostMatrix between unmatched embeddings and embeddings at timestep t .

 - Apply gating to CostMatrix.

 - Apply hungarian to match some of the detections.

- Match the remaining detections with IoU cost matrix (hungarian algorithm).

- Initiate new tracks, delete old ones.

Deep SORT: Questions



Deep SORT: Assignment

Notebook: seminar3-deep-sort.ipynb



Zip archive in Discord



Work in groups




20 minutes



Ask your teammates

Call the lecturer if you need help



Deep SORT: Questions



Metrics

- $\text{FAF}(\downarrow)$: number of false alarms per frame.
- $\text{MT}(\uparrow)$: number of mostly tracked trajectories. I.e. target has the same label for at least 80% of its life span.
- $\text{ML}(\downarrow)$: number of mostly lost trajectories. i.e. target is not tracked for at least 20% of its life span.
- $\text{FP}(\downarrow)$: number of false detections.
- $\text{FN}(\downarrow)$: number of missed detections.
- $\text{ID sw}(\downarrow)$: number of times an ID switches to a different previously tracked object [24].
- $\text{Frag}(\downarrow)$: number of fragmentations where a track is interrupted by miss detection.

Metrics

$$\text{MOTA} = 1 - \frac{\sum_t (\text{FN}_t + \text{FP}_t + \text{IDSW}_t)}{\sum_t \text{GT}_t}$$

- Multiple Object Tracking Accuracy.

$$\text{MOTP} = \frac{\sum_{t,i} d_{t,i}}{\sum_t c_t}$$

- Multiple Object Tracking Precision.
Average overlap between all correctly matched hypotheses and their respective objects.

Datasets and benchmarks

Dataset	Classes	Videos		Avg length (s)	Tracks / video	Min resolution	Ann. fps	Total Eval length (s)
		Eval.	Train					
MOT17 [42]	1	7	7	35.4	112	640x480	30	248
KITTI [25]	2	29	21	12.6	52	1242x375	10	365
UA-DETRAC [64]	4	40	60	56	57.6	960x540	5	2,240
ImageNet-Vid [52]	30	1,314	4,000	10.6	2.4	480x270	~25	13,928
YTVIS [70]	40	645	2,238	4.6	1.7	320x240	5	2,967
TAO (Ours)	833	2,407	500	36.8	5.9	640x480	1	88,605

MOT challenge: <https://motchallenge.net/>

Source: <https://arxiv.org/pdf/2005.10356.pdf>

More trackers

SMOT: Single-Shot Multi Object Tracking:

<https://arxiv.org/pdf/2010.16031v1.pdf>

Fast Online Object Tracking and Segmentation: A Unifying Approach:

<https://arxiv.org/pdf/1812.05050v2.pdf>

FairMOT: On the Fairness of Detection and Re-Identification in Multiple Object Tracking

<https://arxiv.org/pdf/2004.01888v5.pdf>